

Organizational Authenticity and Corporate Value Alignment: A Scalable Text-Based Measurement Pipeline

June 8, 2026

Abstract

This measurement pipeline develops and validates a reproducible text-based approach for studying organizational authenticity as alignment between firms’ stated values and priorities reflected in official disclosures. In a fixed sample of 50 large S&P 500 firms across Technology, Financials, Healthcare, Consumer Discretionary, and Energy from 2016–2024, I collect archived corporate stated-values pages from the Wayback Machine, pair them with SEC definitive proxy statements (DEF 14A), and construct an Organizational Authenticity Index from overlap between 12-theme distributions. The panel retains all 450 target company-years, including 358 usable stated-values observations, 434 collected proxy statements, and 328 scoreable company-years. The mean authenticity index is 41.98 on a 0–100 scale, indicating moderate rather than high alignment between stated values and proxy-disclosure priorities. The proxy corpus is not empty legal boilerplate: shareholder/performance language dominates, but workforce, DEI, leadership, and stakeholder language are pervasive and increasingly salient over time. Robustness checks show that the index is not mechanically driven by proxy-statement length and is only weakly related to whole-text semantic similarity. Measurement diagnostics show that proxy-statement genre pressure is a modest measurement risk rather than the main explanation for low alignment; values evidence is concentrated in governance and meeting-related sections, human-capital sections are most values-dense, and same-theme semantic comparability varies substantially across themes. The analysis is a disclosure-priority alignment proxy, not direct evidence of corporate behavior or ethical authenticity.

1 Introduction

Organizations routinely publish statements about their values, purpose, mission, and commitments. Stakeholders care not only what organizations say they value, but whether those claims align with observable priorities. This study develops a scoped proof of concept for measuring such congruence across a longitudinal sample of large public firms. The empirical question is not whether firms are morally authentic in an absolute sense, but whether themes emphasized on public stated-values pages also appear in official disclosure artifacts that reveal what firms choose, or are required, to discuss.

The analysis uses an integrated measurement pipeline. I collect archived corporate “About,” mission, purpose, and values pages through the Wayback Machine, then collect SEC DEF 14A proxy statements as reproducible public records of disclosed governance and human-capital priorities. I construct an Organizational Authenticity Index from overlap between stated-values and proxy-disclosure theme distributions, then stress-test whether low scores reflect proxy-statement genre, where values evidence appears in proxy text, and whether same-labeled themes are semantically comparable across source types.

The contribution is methodological and empirical. Methodologically, I operationalize an ambiguous social-science construct with transparent collection rules, explicit missingness, auditable dictionary evidence, model-based robustness checks, and validity diagnostics. Empirically, large public firms show moderate alignment between stated-values pages and proxy-disclosure priorities, with meaningful sector and year variation. The index is designed to identify cases for interpretation and audit, not to rank firms as authentically good or bad.

2 Conceptual Framework

I define organizational authenticity as alignment between what organizations say they value and what their behavior suggests they prioritize. Because behavior is not directly observed here, I treat proxy statements as official disclosure evidence. The measurement logic is conservative: values pages and proxy statements are both public artifacts, but they belong to different communication contexts. Alignment between them is evidence of disclosure-priority consistency; divergence is an audit signal requiring interpretation.

This framing follows work on stated-lived value congruence and organizational authenticity, which emphasizes congruence between formal claims and perceived practice [5]. It is also informed by institutional theory, which warns that formal claims may be symbolic or decoupled from organizational activity [4], and by CSR disclosure research showing that external reporting can combine substantive and symbolic elements [7, 1]. The text-analysis strategy follows the domain-specific dictionary caution in financial text research: general sentiment or generic language measures are often inappropriate for regulated corporate disclosures [3].

3 Data

3.1 Sample

The sample consists of 50 firms: the 10 largest S&P 500 firms by market capitalization in each of five sectors: Technology, Financials, Healthcare, Consumer Discretionary, and Energy. The 2016–2024 window yields a balanced target grid of 450 company-years.

3.2 Stated-values source

The stated-values corpus uses official archived corporate pages from the Wayback Machine. Eligible pages include corporate “About Us,” mission, values, purpose, “Who We Are,” or equivalent parent company pages; product, subsidiary, investor-only, campaign, careers, and third-party pages are excluded. For each company-year, the pipeline selects the same-year successful HTML capture nearest June 30 at 12:00 UTC. Adjacent-year substitutions are not used.

All 450 target company-years receive explicit status records. The stated-values corpus contains 358 usable observations (79.6%); the remaining 92 are documented gaps: 60 no CDX capture, 15 retrieval failures, 12 insufficient-substantive-text cases, and 5 no-eligible-capture cases. Usable coverage rises from 29 of 50 firms in 2016 to 46 of 50 in 2024.

3.3 Disclosure source

The disclosure corpus uses SEC definitive proxy statements (DEF 14A) as the lived-values disclosure source. Proxy statements are not direct behavioral evidence, but they are official, free, longitudinally available, and auditable. I use DEF 14A rather than sustainability reports or corporate blogs for

three reasons. First, the filing is available through standardized public infrastructure for nearly all sampled firms. Second, it contains board, compensation, shareholder, human-capital, diversity, risk, and governance discussion, making it a reasonable place to observe priorities firms formally disclose to investors. Third, its legal and governance genre makes it useful for measurement diagnosis: if stated values are absent from this formal channel, that absence is substantively interesting but requires caution. The pipeline resolves tickers to SEC CIKs, retrieves submissions metadata, selects the calendar-year DEF 14A, downloads the primary filing, and extracts clean text with source URLs, accession numbers, hashes, and extraction-quality fields.

The proxy corpus collects 434 of 450 company-years (96.44%). The 16 missing rows are retained with the structured gap reason `no_def14a_filing_for_calendar_year`; missing rows are never interpreted as zero disclosure.

4 Measurement

4.1 Theme taxonomy

Both corpora use a shared deterministic 12-theme taxonomy (Table 1). The taxonomy is deductive rather than outcome-fitted. Its categories combine the stated-values construct with recurring domains in organizational authenticity, CSR, stakeholder, corporate governance, and corporate-culture literatures: stakeholder orientation, employees and human capital, ethics and integrity, DEI, sustainability, shareholder performance, leadership/accountability, collaboration, and organizational purpose [5, 7, 1, 2]. The design is broader than a pure governance dictionary because stated-values pages include aspirational identity claims, but narrower than generic sentiment analysis because regulated disclosures require domain-specific language choices [3].

Table 1: Theme taxonomy for stated-values and proxy-disclosure evidence

Theme no. and label	Construct meaning	Example indicators
1. Customers and service	Customer orientation, service quality, consumer value, and client-facing commitments.	customer, consumer, service, client
2. Employees and workplace	People, workforce, talent, culture, workplace experience, and employee-facing priorities.	employees, workforce, talent, workplace
3. Innovation and excellence	Innovation, quality, improvement, operational excellence, and superior execution.	innovation, quality, excellence, improvement
4. Integrity and ethics	Ethical conduct, trust, transparency, compliance, and responsible business conduct.	integrity, ethics, trust, transparent
5. Diversity, equity, and inclusion	Diversity, equity, inclusion, representation, and inclusive governance or culture.	diversity, equity, inclusion, diverse
6. Social impact and community	Community engagement, social responsibility, philanthropy, and broader societal contribution.	community, society, charitable, social impact
7. Environment and sustainability	Environmental responsibility, sustainability, climate, emissions, and natural-resource stewardship.	sustainability, climate, emissions, environmental
8. Health, safety, and wellbeing	Health, safety, security, wellbeing, and risk prevention for employees or stakeholders.	health, safety, wellbeing, security
9. Shareholders and performance	Shareholder value, financial performance, returns, profitability, and investor-facing priorities.	shareholder, returns, performance, value

Continued on next page

Table 1: Theme taxonomy for stated-values and proxy-disclosure evidence (continued)

Theme no. and label	Construct meaning	Example indicators
10. Leadership and accountability	Leadership, responsibility, accountability, ownership, and governance oversight.	leadership, accountability, responsibility, oversight
11. Collaboration and partnership	Cooperation, partnership, teamwork, stakeholder collaboration, and joint action.	collaboration, partnership, together, teamwork
12. Purpose and identity	Mission, purpose, identity, values, and broad statements about organizational reason-for-being.	purpose, mission, values, identity

Themes are assigned using auditable keyword and phrase evidence. For each positive classification, the pipeline stores matched phrases, excerpts, and counts. This supports reproducibility and audit, but can miss implicit values language or over-count legally repeated terms.

4.2 Authenticity index

For each scoreable company-year, the pipeline converts stated-values and proxy-disclosure evidence into normalized 12-dimensional theme-share vectors. Let s_i be the stated-values share for theme i , and let d_i be the proxy-disclosure share. The primary Organizational Authenticity Index is:

$$A = 100 \times \sum_{i=1}^{12} \min(s_i, d_i). \quad (1)$$

The score ranges from 0 to 100. A score of 100 means the theme distributions are identical; 0 means their nonzero theme emphasis is fully disjoint. Substantively, the score measures the share of stated-values emphasis mirrored in proxy-disclosure emphasis. It is a disclosure-priority alignment proxy, not a direct conduct measure.

This overlap form is appropriate because both inputs are probability distributions over the same theme space. The sum of componentwise minima is the common probability mass; equivalently, it is one minus one-half of the L1 distance between the distributions. The index is bounded, symmetric, interpretable in percentage-point terms, and robust to raw document length after normalization. It also fits the construct: authenticity is operationalized as theme-priority congruence, not absolute volume of values language.

Rows are scored only when both corpora contain usable text and nonzero theme evidence. The analytical panel retains all 450 target company-years: 328 are scored, 92 lack usable stated-values evidence, 16 lack a collected proxy statement, and 14 have insufficient stated-values theme signal.

4.3 Robustness and diagnostic measures

The robustness analysis adds cosine alignment, L1 alignment, and binary Jaccard overlap over the same 12-theme space. These checks test whether the result depends on the overlap functional or is stable across alternative vector-similarity notions. Cosine captures directional similarity in theme emphasis, L1 captures total distributional distance, and Jaccard captures coarse theme co-presence. For stated-values vector $\mathbf{s}$ and disclosure vector $\mathbf{d}$:

$$\text{Cosine}(\mathbf{s}, \mathbf{d}) = \frac{\mathbf{s} \cdot \mathbf{d}}{\|\mathbf{s}\|_2 \|\mathbf{d}\|_2}, \quad (2)$$

$$\text{L1Alignment}(\mathbf{s}, \mathbf{d}) = 1 - \frac{1}{2} \sum_{i=1}^{12} |s_i - d_i|, \quad (3)$$

$$\text{Jaccard}(\mathbf{s}, \mathbf{d}) = \frac{|\{i : s_i > 0\} \cap \{i : d_i > 0\}|}{|\{i : s_i > 0\} \cup \{i : d_i > 0\}|}. \quad (4)$$

Because $\mathbf{s}$ and $\mathbf{d}$ are normalized theme-share vectors, the primary authenticity index is numerically equivalent to $100 \times \text{L1Alignment}$; cosine and Jaccard are reported as conceptually distinct robustness checks.

The expected diagnostic pattern is high agreement among distributional measures and weaker agreement with binary co-presence. If the primary index is stable, it should correlate strongly with cosine and L1 alignment because all three use relative attention across the same 12 themes. Binary Jaccard should be related but weaker because it ignores intensity and only asks whether both texts mention a theme at least once, a low bar for long proxy statements. Conversely, if the index were driven mainly by document length, generic corporate language, or genre boilerplate, it would correlate strongly with proxy word count, whole-text semantic similarity, or proxy-genre pressure.

The analysis also computes supplementary semantic similarity with **sentence-transformers/all-MiniLM-L6-v2** over deterministic representative text windows. Sentence-transformer models are designed for sentence- and passage-level semantic similarity [6], and MiniLM-style distillation provides a lightweight, open-source encoder feasible for reproducible scoring across hundreds of long-document pairs [8]. The embedding score is not primary because broad semantic similarity can reflect corporate genre, legal language, or shared business vocabulary rather than values-theme alignment. It is used as a robustness and contrast measure: agreement with the dictionary index increases stability, while divergence flags audit cases. For embedded text windows $\mathbf{e}^S$ and $\mathbf{e}^D$, semantic similarity is computed as cosine similarity and multiplied by 100 for comparability with the index scale:

$$\text{SemanticSimilarity} = 100 \times \frac{\mathbf{e}^S \cdot \mathbf{e}^D}{\|\mathbf{e}^S\|_2 \|\mathbf{e}^D\|_2}. \quad (5)$$

The measurement-diagnostics layer adds proxy-genre pressure, section-level proxy parsing, and theme-level semantic comparison to evaluate construct validity, not causal effects. Genre pressure tests whether low alignment partly reflects procedural proxy language. Section parsing locates proxy-side evidence. Theme-level semantic comparison asks whether the same taxonomy label carries comparable local meaning across stated-values pages and proxy filings.

5 Results

5.1 Coverage

Table 2 summarizes coverage and scoring. The study preserves all 450 target company-years with explicit statuses rather than silently dropping missing or insufficient observations.

Table 2: Coverage and scoring flow

Status	Company-years	Share of target grid
Target company-years	450	100.0%
Usable stated-values pages	358	79.6%
Collected DEF 14A filings	434	96.4%
Scored authenticity-index rows	328	72.9%
Not scored: no usable stated-values page	92	20.4%
Not scored: no collected proxy statement	16	3.6%
Not scored: insufficient stated-values theme signal	14	3.1%

The main attrition point is the historical stated-values corpus, not SEC proxy coverage. Proxy coverage is 96.4%, while stated-values coverage is lower because some archived pages have no eligible capture, no substantive text, or retrieval failure. Thus unscored rows mostly reflect web-archive limitations rather than zero disclosure.

5.2 Proxy disclosure patterns

Before constructing the alignment index, the proxy-disclosure analysis characterizes the disclosure side of the measurement problem. Normalized by document length, the strongest proxy theme is shareholders and performance, at 37.78 matches per 10,000 words and 100% presence across collected filings, as expected for shareholder-facing governance documents. Workforce and governance-identity language is also pervasive: employees and workplace averages 28.11 matches per 10,000 words, diversity, equity, and inclusion 25.82, and leadership and accountability 24.87; all three appear in every collected company-year.

Table 3: Leading normalized themes in proxy statements

Theme	Mean matches / 10k words	Presence
Shareholders and performance	37.78	100.0%
Employees and workplace	28.11	100.0%
Diversity, equity, and inclusion	25.82	100.0%
Leadership and accountability	24.87	100.0%
Customers and service	16.79	99.8%
Environment and sustainability	13.72	98.8%

The proxy corpus also shows sector and time heterogeneity. Consumer Discretionary, Financials, and Healthcare are especially shareholder/performance-heavy; Technology’s highest normalized theme is DEI at 33.3 matches per 10,000 words; and Energy has unusually high environment/sustainability language at 23.8. In the 2020–2021 window, DEI rises by 6.76 matches per 10,000 words relative to pre-2020, employees/workplace by 6.25, and health/safety by 3.08. Lexical tone also changes: collective voice rises from 1.212 to 1.681 markers per 100 words between 2016 and 2024, while stakeholder orientation peaks in 2021 at 0.324 and remains above its 2016 level in 2024. These findings show systematic values-related disclosure structured by shareholder governance genre.




Figure 1: Sector-theme heatmap for proxy-statement theme emphasis

Figure 1 shows a common shareholder-governance backbone alongside sector-specific vocabularies. The brightest cells are shareholder/performance in Consumer Discretionary (53.0), Healthcare (47.3), and Financials (43.9); Technology is distinctive for DEI (33.3), while Energy is high on employees/workplace (31.3) and environment/sustainability (23.8). Alignment scores should therefore be read with attention to sector-specific disclosure language.




Figure 2: Event-window change in proxy-statement theme emphasis

Figure 2 shows the clearest descriptive increases around 2020–2021 for DEI, employees and workplace, and health/safety language, consistent with a broader shift toward workforce and stakeholder-facing disclosure. The figure is descriptive rather than causal; it shows that the proxy corpus changed in ways the authenticity index must accommodate.

5.3 Authenticity-index distribution

Across 328 scored company-years, the mean authenticity index is 41.98 and the median is 44.83. Scores range from 0.00 to 82.12, with an interquartile range of 30.81–55.14. The typical company-year therefore shows moderate rather than high alignment: for the median row, less than half of stated-values theme emphasis is mirrored in proxy-disclosure theme distribution.




Figure 3: Distribution of the authenticity index

Figure 3 shows the core empirical distribution. Scores are not clustered near either endpoint: the median company-year is moderately aligned, while a sizable lower tail signals sharp divergence between stated values and proxy-disclosure priorities. This supports treating authenticity as a graded alignment construct rather than a binary property.

5.4 Sector and temporal patterns

Sector averages vary descriptively. Energy and Healthcare have the highest average alignment scores (47.03 and 46.74), while Consumer Discretionary and Financials are lower (36.31 and 37.08); Technology sits near the overall mean at 42.58. These patterns should not be read as sector-level behavioral authenticity because they may reflect sector vocabularies, reporting obligations, document genre, and differential stated-values coverage.

Table 4: Authenticity index by sector

Sector	Scored rows	Mean	Median	Range
Consumer Discretionary	59	36.31	41.99	0.51–64.49
Energy	62	47.03	46.03	1.57–82.12
Financials	74	37.08	37.58	0.60–72.50
Healthcare	73	46.74	49.97	6.27–72.76
Technology	60	42.58	45.16	0.00–66.09

Year summaries show modest upward movement after 2019, with average scores rising from 39.68 in 2016 to the low-to-mid 40s from 2019 through 2024. Because scored-row coverage also rises over time, this descriptive pattern may partly reflect improved source availability rather than substantive change.




Figure 4: Mean authenticity index by year

Figure 4 suggests a modest post-2019 increase, but later years also have better stated-values coverage. The year pattern may therefore combine substantive change with improved source availability.

5.5 Robustness checks

Table 5 summarizes robustness and diagnostic checks. The primary index is strongly correlated with cosine alignment ($r = 0.919$) and perfectly correlated with normalized L1 alignment by construction for probability distributions ($r = 1.000$). Binary Jaccard overlap is weaker ($r = 0.778$), confirming that binary theme presence is less informative when proxy statements mention many themes.

Table 5: Robustness and diagnostic check summary

Check	Expected pattern	Observed result	Interpretation
Cosine alignment	Strong positive association with the primary index.	Pearson $r = 0.919$	Theme-priority alignment is stable across distributional similarity metrics.
Normalized L1 alignment	Perfect or near-perfect association by construction for probability vectors.	Pearson $r = 1.000$	Confirms the overlap index is equivalent to common probability mass.
Binary Jaccard overlap	Positive but weaker association because it ignores emphasis intensity.	Pearson $r = 0.778$	Binary co-presence is informative but less precise than theme-share overlap.
Proxy word count	Weak association if length normalization works.	Pearson $r = -0.046$	The index is not mechanically driven by proxy length.
Whole-text semantic score	Low-to-moderate association; embeddings may capture genre rather than values priorities.	Pearson $r = 0.136$ for rows with semantic scores	Semantic similarity is useful as a contrast measure, not a substitute for the theme index.

The near-zero proxy word-count correlation suggests that normalization removes mechanical length effects. Whole-text semantic similarity is available for 342 rows and has low correlation with the primary index among scored rows. This is useful: the embedding measure captures broad rhetorical or genre similarity, whereas the primary index captures relative values-theme emphasis.




Figure 5: Theme-distribution alignment versus semantic similarity

Figure 5 shows why the semantic score is diagnostic rather than a replacement for the index. The four quadrants identify keyword-semantic agreement or divergence; the semantic-high-only quadrant is especially important because generic corporate or legal language may mask theme-priority divergence.




Figure 6: Distribution of robustness metrics

Figure 6 compares keyword theme-overlap, whole-text semantic similarity, and an equal-weight hybrid score. The semantic distribution is more compressed and centered higher than the keyword index, while the hybrid falls between them. This supports keeping the keyword index primary and using semantic similarity as a contrast diagnostic.

6 Exploratory Measurement Diagnostics

The diagnostic analysis asks whether proxy-statement genre and theme semantics affect the index. It addresses three questions: (1) Is the index sensitive to whole-document proxy genre pressure? (2) Which proxy sections carry values-theme evidence? (3) Are same-labeled themes semantically comparable across stated-values pages and proxy statements?

6.1 Proxy genre pressure

Proxy genre pressure is computed from three phrase families: shareholder-meeting mechanics, governance boilerplate, and legal/procedural language. Each count is normalized per 1,000 words, z-scored, and averaged. Standardization prevents the most frequent phrase family from mechanically dominating the composite while treating the three component rates as conceptually comparable. Higher values indicate stronger procedural and governance-genre language. The composite is:

$$\text{GenrePressure}_{cy} = \frac{1}{3} [z(M_{cy}) + z(G_{cy}) + z(L_{cy})], \quad (6)$$

where c indexes company and y year. M_{cy} is the shareholder-mechanics rate, G_{cy} the governance-boilerplate rate, and L_{cy} the legal/procedural rate. Each component is phrase matches per 1,000 proxy words, and $z(\cdot)$ standardizes among collected proxy rows. Proxy genre pressure is weakly negatively correlated with the authenticity index: Pearson $r = -0.093$ and Spearman $\rho = -0.141$. Mean authenticity is 44.19 in the low genre-pressure tercile and 39.70 in the high tercile, a 4.48-point difference.

A descriptive fixed-effects regression with proxy word count, sector indicators, and year indicators yields a proxy-genre-pressure coefficient of -2.764 , with $R^2 = 0.096$. This is not a causal estimate: controls absorb basic document scale, sector vocabulary, and period effects rather than identify an exogenous genre shock. Genre pressure is therefore a modest measurement-risk signal, not the main explanation for low scores.




Figure 7: Proxy genre pressure and authenticity

Figure 7 shows diffuse scatter rather than a strong visual relationship, matching the diagnostics: Pearson $r = -0.093$ is not conventionally significant ($p = 0.093$), while Spearman $\rho = -0.141$ is significant but small ($p = 0.011$). Genre pressure is a modest measurement concern, not the main explanation for low alignment.

6.2 Section-level proxy parsing

The section parser identifies 127,630 section-like spans across 434 collected proxy statements and assigns them to governance/board, meeting/voting, compensation, shareholder proposals, ownership, audit, legal/other, human-capital/values, and related families. This heuristic structural parser uses

headings and section cues to locate values language; it is not a legal-document parser. Its construct-validity purpose is to distinguish evidence embedded in voting instructions or table-of-contents fragments from evidence in governance, human-capital, compensation, and sustainability contexts. Governance/board sections are the largest raw source, with 129,840 theme matches; meeting/voting sections contribute 67,314, and human-capital/values sections 59,405.

Raw volume and density tell different stories. Governance/board sections dominate raw evidence because they are large: 7,523,017 parsed words. Human-capital/values sections are most values-dense, with 36.91 theme matches per 1,000 section words. Thus a theme match in a governance section may carry a different textual meaning than the same match in a human-capital or values section.




Figure 8: Section size and values-theme density

Figure 8 visualizes the size-density distinction: the x-axis is total parsed section volume, the y-axis values-theme density, and bubble size total theme matches. Governance/board sections produce the largest raw volume, while human-capital/values sections are far more values-dense. Raw counts alone would overstate governance machinery and understate direct people-oriented values evidence.




Figure 9: Theme composition by proxy section family

Figure 9 should be read row-wise: each cell is a within-section-family theme share, so brighter cells mark dominant themes within that family. The highlighted cells separate proxy machinery from

more direct values evidence: ownership sections concentrate on leadership/accountability (0.54), meeting/voting and shareholder-proposal sections on shareholder/performance (0.43 and 0.31), and compensation and human-capital/values sections on employees/workplace (0.32 and 0.27).

6.3 Theme-level semantic comparison

The theme-level diagnostic compares stated-values and proxy-disclosure excerpts within the same taxonomy theme. MiniLM semantic similarity is computed for 2,099 company-year-theme pairs where both sides have evidence. This is stricter than whole-text semantic similarity because it asks whether local language attached to the same theme label is similar. It addresses a dictionary validity concern: two texts can match the same theme label while using it differently. For example, “employees” on a values page may refer to culture and belonging, while “employees” in a proxy may appear in compensation or risk governance language.

Theme-level comparability varies sharply. Innovation and excellence has the highest mean local semantic similarity (46.19), followed by environment and sustainability (45.83) and purpose and identity (44.90). Employees and workplace is lowest (25.96), suggesting that stated-values pages and proxy statements discuss employees in different local contexts: culture and people on the stated-values side versus compensation, governance, workforce-risk, or human-capital disclosure in proxy text. Customers and service has the largest mean stated/disclosure emphasis gap (0.167).




Figure 10: Theme semantic similarity and emphasis gap

Figure 10 plots each taxonomy theme by its mean absolute stated/proxy theme-share gap on the x-axis and mean local semantic similarity on the y-axis; bubble size is the number of computed company-year-theme pairs. Customers/service has the largest emphasis gap despite moderate semantic similarity, while employees/workplace is the clearest low-similarity outlier. The figure shows why theme-level diagnostics add information beyond aggregate alignment.

6.4 Illustrative case-audit targets

The aggregate index is designed to produce audit targets rather than final judgments. Table 6 reports four illustrative company-years selected by the two primary diagnostic dimensions: Organizational

Authenticity Index (OAI) and whole-text semantic similarity (SS). Proxy genre pressure (GP) is retained only as auxiliary context, not as a case-type criterion.

Table 6: Illustrative diagnostic case-audit targets

Case type	Company-year	OAI	SS	GP	Audit interpretation
Low OAI, low SS	NVIDIA 2021	0.00	51.46	-0.51	Low alignment on both measures.
Low OAI, high SS	Goldman Sachs 2020	9.02	70.76	0.92	Broad language similarity but weak theme alignment.
High OAI, low SS	AbbVie 2020	72.76	50.96	0.23	Theme alignment with weaker whole-text similarity.
High OAI, high SS	Valero Energy 2024	81.28	67.81	-0.46	Robust alignment across both measures.

This table is a triage device, not a qualitative validation exercise. The same index value should be interpreted with diagnostic context. GP should be read after the OAI/SS pattern: higher GP may indicate proxy mechanics diluting values evidence, while lower GP makes substantive stated/proxy theme mismatch more plausible.

7 Discussion

The findings support three conclusions. First, large public firms show moderate, not high, disclosure-priority alignment between stated-values pages and proxy statements. The mean authenticity index is 41.98, the median is 44.83, and sector averages range from 36.31 in Consumer Discretionary to 47.03 in Energy. These central results turn two public textual corpora into a comparable company-year measure. Second, the index is not reducible to proxy boilerplate: proxy genre pressure has a weak negative relationship with alignment but does not explain most variation. Third, the source of values evidence matters. Proxy-side theme matches often arise in governance and meeting machinery, while human-capital/values sections are more values-dense. Same theme labels are not always semantically equivalent across source types; employees/workplace is common in both sources but differs substantially in local meaning.

These results support interpreting the index as an audit tool. High scores indicate public-artifact consistency across two communication channels. Low scores indicate divergence that may reflect true differences in emphasis, source missingness, narrow stated-values pages, document-genre effects, or taxonomy limitations. The diagnostics help distinguish these explanations.

8 Limitations

The most important limitation is construct validity. Proxy statements are official disclosure documents, not direct observations of organizational behavior, so the index measures disclosure-priority alignment rather than culture or ethical conduct. A second limitation is model dependence: the deterministic taxonomy is auditable but can miss synonyms or context; semantic checks help but do not solve this issue. Third, missingness is nonrandom because only 328 of 450 company-years can be scored. Fourth, the heuristic section parser can misclassify filing fragments, especially when extracted proxy text contains table-of-contents artifacts or repeated headings. Finally, all external-event interpretations are descriptive and should not be read causally.

9 Conclusion

This study demonstrates a scalable design for studying organizational authenticity as alignment between stated values and official disclosure priorities. It collects historical stated-values pages, pairs them with reproducible proxy disclosures, constructs an auditable theme-overlap index, and stress-tests the index with semantic and genre diagnostics. The pipeline preserves the full 450 company-year sampling frame, documents missingness, and produces 328 scoreable observations, making coverage part of the measurement object rather than an unreported preprocessing loss.

The main empirical result is moderate, not high, alignment. The mean Organizational Authenticity Index is 41.98 on a 0–100 scale, with meaningful sector and year variation. Technology has the highest sector mean, Consumer Discretionary the lowest, and average alignment rises after 2020. The proxy corpus also contains substantive values evidence rather than only legal template language: shareholder/performance themes dominate, but employee, DEI, leadership, environmental, and stakeholder themes appear throughout the corpus and vary by sector and period.

The diagnostic evidence clarifies what the index should and should not be asked to do. Robustness checks show that the index is not mechanically driven by proxy length or by whole-text semantic similarity, and regression diagnostics indicate that proxy genre pressure is a modest measurement risk rather than the main explanation for low scores. Section parsing shows that values language is not evenly distributed across proxy statements: governance and meeting sections contain large volumes of values-relevant language, while human-capital sections are smaller but more values-dense. Theme-level semantic comparison further shows that some categories travel well across stated-values pages and proxy disclosures, whereas employee/workplace language is especially context-dependent.

The resulting measure is therefore not a verdict on corporate authenticity, ethical behavior, or organizational quality. It is a transparent instrument for identifying where public values claims and official disclosure priorities align, diverge, or require qualitative review. Its strongest use is as a reproducible screening and diagnostic tool: it can surface sector patterns, temporal shifts, metric disagreements, and case-audit targets, while leaving final interpretation to contextual reading and future validation against behavioral evidence.

References

- [1] Hawn, O., and Ioannou, I. 2016. “Mind the Gap: The Interplay Between External and Internal Actions in the Case of Corporate Social Responsibility.” *Strategic Management Journal* 37(13): 2569–2588.
- [2] Guiso, L., Sapienza, P., and Zingales, L. 2015. “The Value of Corporate Culture.” *Journal of Financial Economics* 117(1): 60–76.
- [3] Loughran, T., and McDonald, B. 2011. “When Is a Liability Not a Liability? Textual Analysis, Dictionaries, and 10-Ks.” *Journal of Finance* 66(1): 35–65.
- [4] Meyer, J. W., and Rowan, B. 1977. “Institutionalized Organizations: Formal Structure as Myth and Ceremony.” *American Journal of Sociology* 83(2): 340–363.
- [5] Pamphile, V. D., and Ruttan, R. L. 2023. “The (Bounded) Role of Stated-Lived Value Congruence and Authenticity in Employee Evaluations of Organizations.” *Organization Science* 34(6): 2332–2351.

- [6] Reimers, N., and Gurevych, I. 2019. “Sentence-BERT: Sentence Embeddings Using Siamese BERT-Networks.” In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*.
- [7] Shabana, K. M., and Ravlin, E. C. 2016. “Corporate Social Responsibility Reporting as Substantive and Symbolic Behavior: A Multilevel Theoretical Analysis.” *Business and Society Review* 121(2): 297–327.
- [8] Wang, W., Wei, F., Dong, L., Bao, H., Yang, N., and Zhou, M. 2020. “MiniLM: Deep Self-Attention Distillation for Task-Agnostic Compression of Pre-Trained Transformers.” In *Advances in Neural Information Processing Systems*.